

Created by the Aerospace Cyber Defence Centre of Excellence (CEC)

Validation report

WAC (Windows Artefacts Collector)

Windows artefact collection software
for digital forensic investigation



SYSTEM
ARTEFACTS



USER
FILES



DIGITAL
EVIDENCE

A TOOL IN SUPPORT
OF DIGITAL
INVESTIGATION

WAC — Windows Artefacts Collector

[Validation report](#)

Software version 1.3.1

Document edition September 2026

Publisher Aerospace Cyber Defence Centre of Excellence (CEC)
French Air and Space Force Academy, Air Base 701, Salon-de-Provence

This document lists the checks that establish the quality of WAC: those WAC runs itself during every collection, the test programs that confront its decoders with a reference, and the automatic checks of its outputs. For each, it states what is compared, against what, the reference result and the defects it has already caught. It also states what these checks do not guarantee. How to run each test is described in the *Testing guide*.

Contents

1	The approach	2
1.1	Why valid JSON proves nothing	2
1.2	Four levels of checks	2
1.3	The development rules that go with them	2
2	The checks built into WAC	4
2.1	Before any extraction	4
2.2	Raw reading of the NTFS volume	4
2.3	Registry hives	4
2.4	Exhibit integrity: the exhibit store	5
2.5	Authenticity of binaries	5
2.6	Event logs	5
2.7	Decoding the artefacts	6
3	The test programs	7
3.1	Format harnesses	7
3.2	Conversions: system_conversions_test in detail	8
3.3	Robustness harnesses	8
4	The automatic cycle and its checks	10
4.1	The cycle	10
4.2	The checks of check-json.py	10
4.3	Reference result	12
5	Cross-checks on a real collection	13
6	What these checks guarantee	14
6.1	Guarantees	14
6.2	Limits	14

The approach

1.1 Why valid JSON proves nothing

WAC produces JSON files meant for a forensic report. Their validity says almost nothing about their accuracy: a field read at the wrong offset, a date interpreted in the wrong time zone, a name decoded in the wrong code page, a wrong decompression all give perfectly valid JSON, and wrong. The most serious defects found in WAC were all of that kind. None showed at run time; each only appeared when a value was confronted with an **independent source**.

WAC's validation therefore rests on one principle: every piece of data that can be cross-checked is, by a means that shares nothing with the code under test.

1.2 Four levels of checks

The checks built into WAC (chapter 2): run during every collection, on the examined machine. They refuse inconsistent data rather than publish it, and attest the integrity of the exhibits taken.

The test programs (chapter 3): fourteen executables that confront each decoder with a reference (the Microsoft function it replaces, a published standard, another implementation), or that prove a parser never reads outside its buffer.

The automatic cycle and its checks (chapter 4): a full collection in a Windows 11 virtual machine, then `check-json.py`, which checks the validity of the outputs and confronts them with one another and with external sources.

Cross-checks on a real collection (chapter 5): what a fresh virtual machine does not contain (old logs, fragmented files, exotic names) is checked on the data of a real machine, processed locally.

1.3 The development rules that go with them

1. **A fix comes with the test that would have caught the defect**, and that test is checked both ways: it passes on the fixed code, and it fails when the defect is reintroduced. A test never seen failing proves nothing.
2. **The build is warning-free**, with `-Wall -Wextra` enabled. The warnings also served as a to-do list: each one marked a field decoded then dropped.

3. **After any change, the new collection is compared with the previous one**, artefact by artefact: an artefact gone from 1,522 entries to 573 lights no warning, only the comparison shows it.
4. **Every input is hostile**, including a file of the examined disk: every read is bounded by the real size of the buffer, never by a length announced in the data.
5. **A field without a value is not emitted**: an empty or null key would read as data found, or as a failed read.

The checks built into WAC

These checks are part of the collection program. They run during every collection, and their result is recorded in `investigation.json`, operation by operation, with its return code and the trace it leaves on the examined system.

2.1 Before any extraction

- Windows version (10 or later) and effective administrator rights (elevated token);
- usable collection location and free space, checked *before* the first extraction: a collection stopped for lack of space would leave an incomplete exhibit store.

2.2 Raw reading of the NTFS volume

Locked files (hives, logs) are read directly from the volume, bypassing the Windows file system. Every structure read is checked before use:

- FILE signature of every \$MFT record and INDX signature of every index block, then check of the *fixups* (update sequence array): every 512-byte stride must end with the record's sequence number. A record *torn* by an interrupted write, a mix of two versions, is refused instead of being read as plausible data;
- LZNT1 and WOF decompression bounded by the expected size of each unit: a unit that does not produce the expected number of bytes is an error, not data silently truncated;
- extracted size compared with the size declared by the \$DATA attribute: their difference, published in the manifest, reveals a truncated extraction a fingerprint alone would not show.

2.3 Registry hives

WAC reads the hives with its own reader, linked into the executable, and not with a library of the examined machine.

- base block: regf signature, checksum (exclusive OR of the first 127 words), primary and secondary sequence numbers. A hive whose logs have not been replayed is **refused**: read as it is, it would give stale keys without any error;

- replay of the transaction logs (.LOG1/.LOG2): signature and sequence chaining of every entry; an entry out of the chain is discarded and counted; an undo journal is written before any change, so that the raw copy can be rebuilt;
- every offset read in the hive is checked against the file's real size, and a subkey pointing to one of its ancestors is refused: a walk cannot loop on a forged hive.

2.4 Exhibit integrity: the exhibit store

Every extracted exhibit is placed in an *exhibit store* that is never reopened for writing; all the work is done on a second copy.

- three fingerprints per exhibit (MD5, SHA-1, SHA-256), computed while reading the volume;
- the working copy is checked by its SHA-256 against what was read from the volume, not against a re-read of the exhibit store;
- the manifest identifying the exhibits is sealed by a second file carrying its fingerprint: a retouched manifest is detected;
- a failed extraction is recorded too: an exhibit missing from the manifest would read as never looked for.

A third party can check the exhibit store without WAC, with the scripts `verify-exhibits.ps1` (English messages) and `verifier-consigne.ps1` (French messages), shipped with the user documentation: they check the seal, then the fingerprint of every exhibit.

2.5 Authenticity of binaries

With the `--binary` option, WAC takes the binaries cited by the artefacts, except those it establishes as signed by Microsoft. That decision rests on a complete verification, done by WAC itself: RSA signature of each security catalog, certificate chain up to a Microsoft root, Authenticode digest of the file found in a verified catalog, or embedded signature. A file whose authenticity is not established is collected: doubt leads to collecting, never to discarding.

2.6 Event logs

Signature of the file, of each chunk and of each record. Every physical chunk of the file is walked, not only those the header declares: a log not properly closed thus yields the records it does not count. An unreadable chunk is reported and skipped; it no longer stops the reading of the file.

2.7 Decoding the artefacts

- every read is bounded by the size of the buffer; an unreasonable number of elements in a list of values (beyond 65,536) is refused rather than walked;
- an impossible date (month 13, hour 25) or one outside the FILETIME range gives an absent date, hence not emitted, and never an arbitrary value;
- every local date carries the offset in force *at its date*, according to the daylight saving rules of the examined machine, year by year.

The test programs

Fourteen executables, left out of WAC's build by their `_test.cpp` suffix and built by `build-windows.sh --test`. Each targets one component and confronts it with a judge that shares nothing with it.

3.1 Format harnesses

Program	Judge	Reference result
sha_test	FIPS 180-4 vectors, independent implementation	SHA-1 and SHA-256 conform, including lengths 55 to 128 bytes, which exercise the padding
lznt1_test	Microsoft's compressor (RtlCompressBuffer)	16 compressed units out of 16 conform, 1,048,576 bytes identical
xpress_test	Microsoft's compressor	137 compressed chunks out of 137 conform, byte for byte
wevt_test	files of a real event provider	the whole message production chain checked, substitution marks included
event_messages_test	SOFTWARE hive and binaries of the machine	the same chain, step by step, starting from the hive
evtx_test	python-evtx, independent implementation	record by record; whole chain identical when streamed and in memory (26 MB peak against 1,281 MB)
offline_registry_test	Microsoft's offreg.dll	18 hives, ten of them from a real machine: 1,273,652 keys and 2,456,139 values identical
system_conversions_test	the Windows functions replaced	9,661,288 comparisons, no difference (details below)
consigne_test	the exhibit store itself, before and after	exhibit store byte-identical after the working copy was modified; seal conforms

Program	Judge	Reference result
ntfs_fixup_test	hand-built records, independent check	sound records restored to the byte, torn records and malformed arrays refused, 20,000 random records against a forbidden page: 44,610 checks, no failure
raw_hive_test	Windows's reg load	the extracted hive is refused before repair and accepted after: the test fails both ways if it must
authenticode_test	Get-Authenticode Signature	2,233 binaries: 2,126 authenticated by both, 99 collected by both, no file accepted by WAC and rejected by Windows; 462 PowerShell and 11 WSH scripts conform

3.2 Conversions: system_conversions_test in detail

WAC depends on five system libraries only (ADVAPI32, KERNEL32, msvcrt, Secur32, WTSAPI32), needed to query the running system. The conversions it used to ask of other libraries of the examined machine are done by its own code, and this program confronts them with the Windows functions they replace:

- GUIDs, SIDs both ways, OLE dates, property names (a table of 1,698 entries generated from the Windows schema);
- FAT dates, on every possible date;
- daylight saving time: offset both ways on the **141 time zones** of Windows, with their rules year by year, at every transition from 1980 to 2035 (hour skipped, hour repeated, change of rule on 1 January);
- reading of the TimeZoneInformation key, compared with GetTimeZoneInformation;
- time zone rules forged at random: no overflow.

Intended, documented differences: WAC refuses a SID text Windows would truncate, and a non-numeric OLE date for which Windows returns a meaningless time.

3.3 Robustness harnesses

A parser reads at offsets the file itself declares. On a truncated or forged file, a read outside the buffer does not crash on its own: it reads the neighbouring memory and publishes it, in valid JSON. The robustness harnesses make it crash: each input is placed against a forbidden memory page, once at its end, once at its start. A single byte read outside the input stops the test. A watchdog also stops a parser that does not end.

Each file is parsed whole, in its truncations, and in randomly corrupted copies (fixed seed: a failure reproduces identically).

Program	Files of a real machine	Inputs parsed
lnk_test	220 shortcuts	546,214
parsers_test	48 automatic jump lists	175,585
	23 custom jump lists	79,036
	307 Prefetch files	1,246,491
	4 registry hives	16,207

Result: no read outside the buffer, no endless loop.

The automatic cycle and its checks

4.1 The cycle

`vmtest/run-wac-test.sh` runs, with no interaction at all, a full collection in a Windows 11 virtual machine (real NTFS, real registry):

- 1. with `--build`, warning-free build of WAC and of the test programs;
- 2. validation of the raw extraction of a hive: `reg load` must refuse it before repair and accept it after;
- 3. run of WAC with SYSTEM rights; the exit code is checked, and the trace is searched for any unhandled exception;
- 4. retrieval of the outputs, of the collection log and of the list of files actually present in the exhibit store;
- 5. run of `check-json.py`.

4.2 The checks of `check-json.py`

The script first checks the JSON validity of every file and the form of the Windows paths (a double escaping shows after parsing). It then confronts values computed by WAC with independent values: another piece of data of the same collection, or an external source. Each check below was written after a real defect, which it caught.

Check	Defect it caught
<i>Dates</i>	
<code>X / XUtc</code> pairs, lists included, naming the same instant	double time-zone shift (sessions, shortcuts, installation date); Prefetch run times labelled <code>+02:00</code> instead of UTC (1,273 values 2 h off)
offset of every local date / IANA time zone database	every date labelled with the collection day's offset (106 winter dates at <code>+02:00</code>); daylight saving dates of the SYSTEM hive read in the wrong layout
no session earlier than the system boot	a two-hour shift on the sessions; later, a boot time 3.5 s late
boot time / Kernel-General 12 event	two independent sources of the same instant: 0.00 s apart

Check	Defect it caught
ShimCache date / NTFS date of the file in the manifest	a UTC date treated as local
no event later than the collection	shift or misreading of an event date
FAT dates without a fraction and with even seconds; Amcache text dates without a fraction	every date written with seven digits of fraction, claiming a precision the source does not have (808 dates)
dates of BAM, UserAssist, USBSTOR and Amcache / taskkill run by the harness, Prefetch, Get-PnpDeviceProperty, PE headers	these four artefacts store UTC, read as local time: all their dates 2 h off; BAM and UserAssist emptied by a regression of the value reading
<i>Decoding</i>	
identifier and serial number of a USB key / Get-PnpDevice	no key carried its serial number, vendor or product: they are in the key names, not in values
Prefetch hash / suffix of its file name	65 wrong hashes out of 270 (hexadecimal without padding)
Prefetch \VOLUME{guid} paths resolved	20,066 paths out of 20,066 unresolved
DestList entries / jump list shortcuts	the first 15 entries of every list lost (573 shortcuts published out of 1,522)
machine name of the events / OperatingSystem.json	check of the BinXML decoding, on two unrelated sources
record identifiers unique per log	a stale chunk read again
unresolved %nnnn references in messages	a provider's parameter file not loaded (3,780 messages)
RID found at the end of the rebuilt SID	reading of the SAM database
service states and types, presence of drivers	enumeration filters mistaken for states
no process carries WAC.exe among its modules	the tool's modules attributed to the Idle process
plausible \$MFT references	decoding of the shortcuts' extension blocks
mounts / Windows' reading of the key (reg query) and partitions (Get-Partition)	6 mounts out of 7 published as "\ (device paths decoded as ANSI)
<i>Exhibit store and replay</i>	
seal / real fingerprint of the manifest	a retouched manifest
three fingerprints per exhibit, consistent counts	an unidentified exhibit

Check	Defect it caught
files present in the exhibit store / manifest	209 exhibits added after sealing, which every other check let through
no hive both replayed and repaired	a replay that does not align the sequences
an undo journal for every replay	a raw copy impossible to rebuild

4.3 Reference result

Last cycle, on version 1.3.1 (Windows 11, Paris time zone):

- 24 valid JSON files out of 24 produced, USBSTOR included: the harness attaches a virtual USB key on every cycle, so that a recreated test machine knows one too;
- 33 cross-checks passed, no failure, among them 4,068 local dates confronted with the IANA database and 145,645 events decoded;
- 1,336 exhibits in the exhibit store, exactly those of the manifest, seal conforms.

Cross-checks on a real collection

A fresh virtual machine produces neither fragmentation, nor old jump lists, nor exotic file names, nor damaged logs. Several defects only showed on a real machine, confronted with an independent source:

Cross-check	What it showed
jump lists / olefile	the streams are named "1" to "f", not "01" to "0f"
shortcut target / DestList entry	a wrong code page and an ignored Unicode path
logs / python-evtx	a chunk with a wrong signature stopped the reading of the whole file (14 records read out of 270)
EVTX chunks / their CRC32	one-off check: the 1,505 chunks of the 404 logs of the test machine are intact
hives / offreg.dll	1,273,652 identical keys
binaries / Get-AuthenticodeSignature	no undue authentication
old output / new one	the same files decoded by the current code, compared after normalisation

Keep in mind

A real collection contains personal data. It is processed locally; a temporary copy in the virtual machine or in a working folder is deleted after use, and nothing of it enters the source repository.

What these checks guarantee

6.1 Guarantees

- **Exhibit integrity:** every exhibit is identified by three fingerprints, the working copy is checked against the reading of the volume, the manifest is sealed, and the exhibit store is never modified (`consigne_test`, checks of the cycle).
- **No writing to the examined system,** apart from the unavoidable traces of a collection on a running system, all recorded in `investigation.json` with the expected trace.
- **Accuracy of the decoders:** where a reference exists — hive reader, decompressions, fingerprints, conversions, time zones, authenticity of binaries — results identical to it on the volumes stated.
- **Robustness:** no shortcut, jump list, Prefetch or hive parser reads outside its buffer or loops, on more than two million truncated or corrupted inputs.
- **Consistency of the outputs:** every possible cross-check between two artefacts, or between an artefact and an external source, is run in every cycle.
- **Independence from the examined machine:** WAC borrows no library from it to interpret the evidence.

6.2 Limits

Limit

These checks do not prove the absence of defects. They prove that the defects they target are absent, on the data they have seen.

- **Completeness:** an empty artefact cannot be told apart from a machine without traces. Only the comparison between two collections, or with a machine whose content is known, measures it.
- **Data the test machine does not hold:** MTP devices, files spread over several \$MFT records (\$ATTRIBUTE_LIST), very large volumes. They are to be validated on real systems.
- **Artefacts without a reference:** when no independent tool reads a format, its accuracy rests on the internal cross-checks and on the public documentation of the format.
- **Time zones:** the confrontation with the IANA database covers the zones whose rules agree with Windows's over the period checked (Western Europe, United States, UTC);

the others are checked against Windows only.

- **MSVC build:** the Visual Studio project is maintained, but the checks are made on the MinGW build, the one delivered.